

Introduction to Data Science

Final Project

Topic :
RCNN-Family

111550038	陳易
111550151	徐嘉亨
111550177	吳定霖

R-CNN

1. Region Proposals (selective search)
2. Feature Extractor (AlexNet)
3. Detector (SVM classifier, bbox regressor)

Region Proposals (selective search)



Image from Jasper RR Uijlings [1]

Region Proposals (selective search)

Algorithm 1: Hierarchical Grouping Algorithm

Input: (colour) image

Output: Set of object location hypotheses L

Obtain initial regions $R = \{r_1, \dots, r_n\}$ using [13]

Initialise similarity set $S = \emptyset$

foreach *Neighbouring region pair* (r_i, r_j) **do**

 Calculate similarity $s(r_i, r_j)$

$S = S \cup s(r_i, r_j)$

while $S \neq \emptyset$ **do**

 Get highest similarity $s(r_i, r_j) = \max(S)$

 Merge corresponding regions $r_t = r_i \cup r_j$

 Remove similarities regarding $r_i : S = S \setminus s(r_i, r_*)$

 Remove similarities regarding $r_j : S = S \setminus s(r_*, r_j)$

 Calculate similarity set S_t between r_t and its neighbours

$S = S \cup S_t$

$R = R \cup r_t$

Extract object location boxes L from all regions in R

Feature Extractor (AlexNet)





Figure 2: An illustration of the architecture of our CNN, explicitly showing the delineation of responsibilities between the two GPUs. One GPU runs the layer-parts at the top of the figure while the other runs the layer-parts at the bottom. The GPUs communicate only at certain layers. The network's input is 150,528-dimensional, and the number of neurons in the network's remaining layers is given by 253,440–186,624–64,896–64,896–43,264–4096–4096–1000.

Feature Extractor (AlexNet)

1. Replacing Sigmoid with ReLU in LeNet
2. Replacing Avg Pooling with Max Pooling
3. Use Dropout in sixth and seventh to replace Fully Connected NN

Feature Extractor (AlexNet)



Detector (SVM classifier, Bbox regressor)

1. SVM classifier -> classification
2. Bbox regressor -> regression

Detector (SVM classifier, Bbox regressor)

$$\hat{G}_x = P_w d_x(P) + P_x$$

$$\hat{G}_y = P_h d_y(P) + P_y$$

$$\hat{G}_w = P_w \exp(d_w(P))$$

$$\hat{G}_h = P_h \exp(d_h(P)).$$

$$t_x = (G_x - P_x)/P_w$$

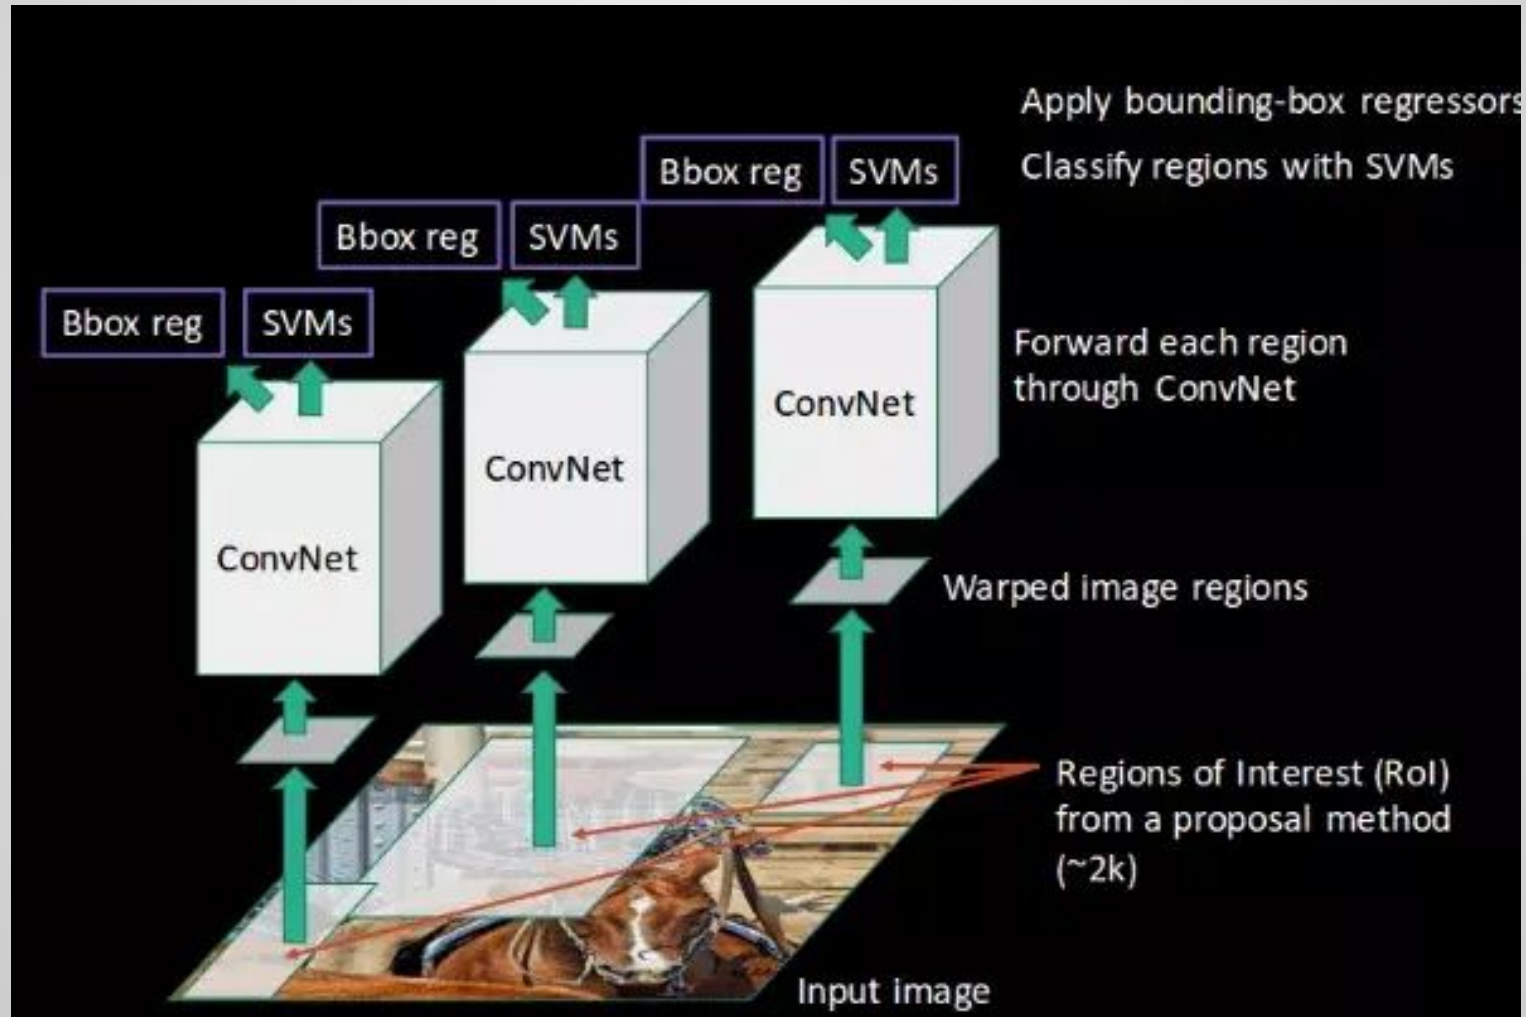
$$t_y = (G_y - P_y)/P_h$$

$$t_w = \log(G_w/P_w)$$

$$t_h = \log(G_h/P_h).$$

$$\mathbf{w}_\star = \underset{\hat{\mathbf{w}}_\star}{\operatorname{argmin}} \sum_i^N (t_\star^i - \hat{\mathbf{w}}_\star^\top \phi_5(P^i))^2 + \lambda \|\hat{\mathbf{w}}_\star\|^2$$

Detector (SVM classifier, Bbox regressor)





Shortcoming

1. Need large memory to store feature vector of proposal
2. Many times convolution training
3. 47 seconds per detection

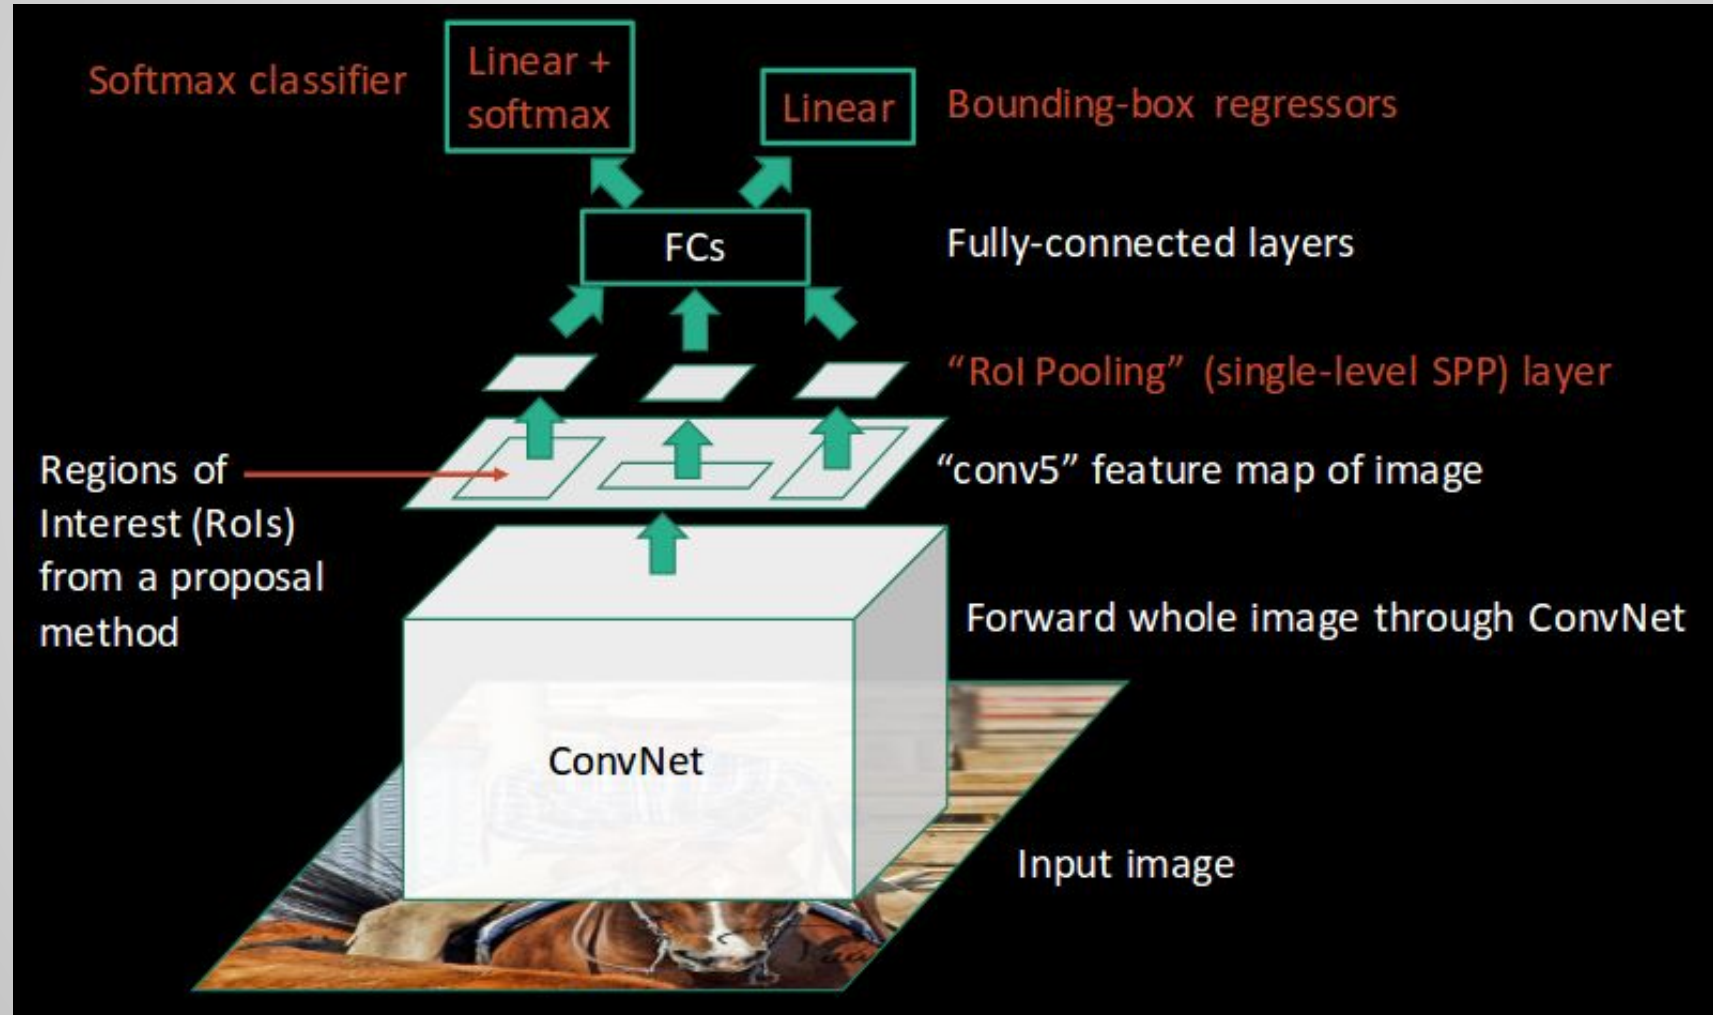
Fast R-CNN

1. Pass the entire image through convolution to generate a feature map.
2. Utilize ROI pooling layer to standardize dimensions.
3. Replace SVM with softmax for classification in the neural network.


Region of Interest (selective search)




Region Proposals (selective search)



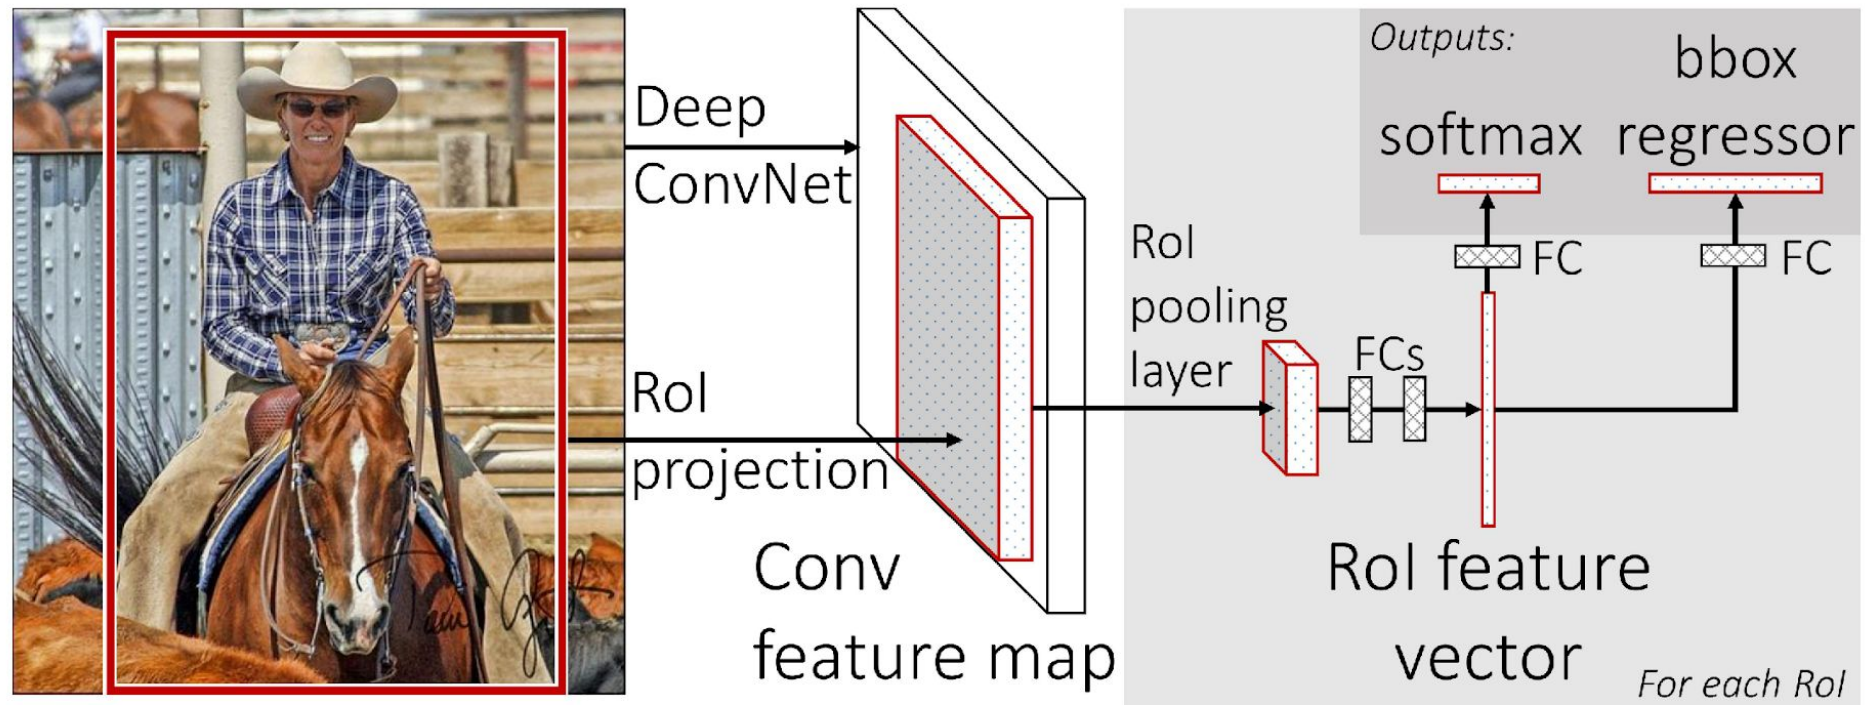
Feature Extractor (VGG net)



ROI Pooling Layer



Detector



Loss Function

$$L(p, u, t^u, v) = L_{\text{cls}}(p, u) + \lambda[u \geq 1]L_{\text{loc}}(t^u, v), \quad (1)$$

$$L_{\text{loc}}(t^u, v) = \sum_{i \in \{x, y, w, h\}} \text{smooth}_{L_1}(t_i^u - v_i), \quad (2)$$

in which

$$\text{smooth}_{L_1}(x) = \begin{cases} 0.5x^2 & \text{if } |x| < 1 \\ |x| - 0.5 & \text{otherwise,} \end{cases} \quad (3)$$

Shortcoming

1. 25 times faster, but still too slow
2. Region Proposal bottle neck

Faster R-CNN

1. Use Anchor
2. Replace Selective Search with Region Proposal Network


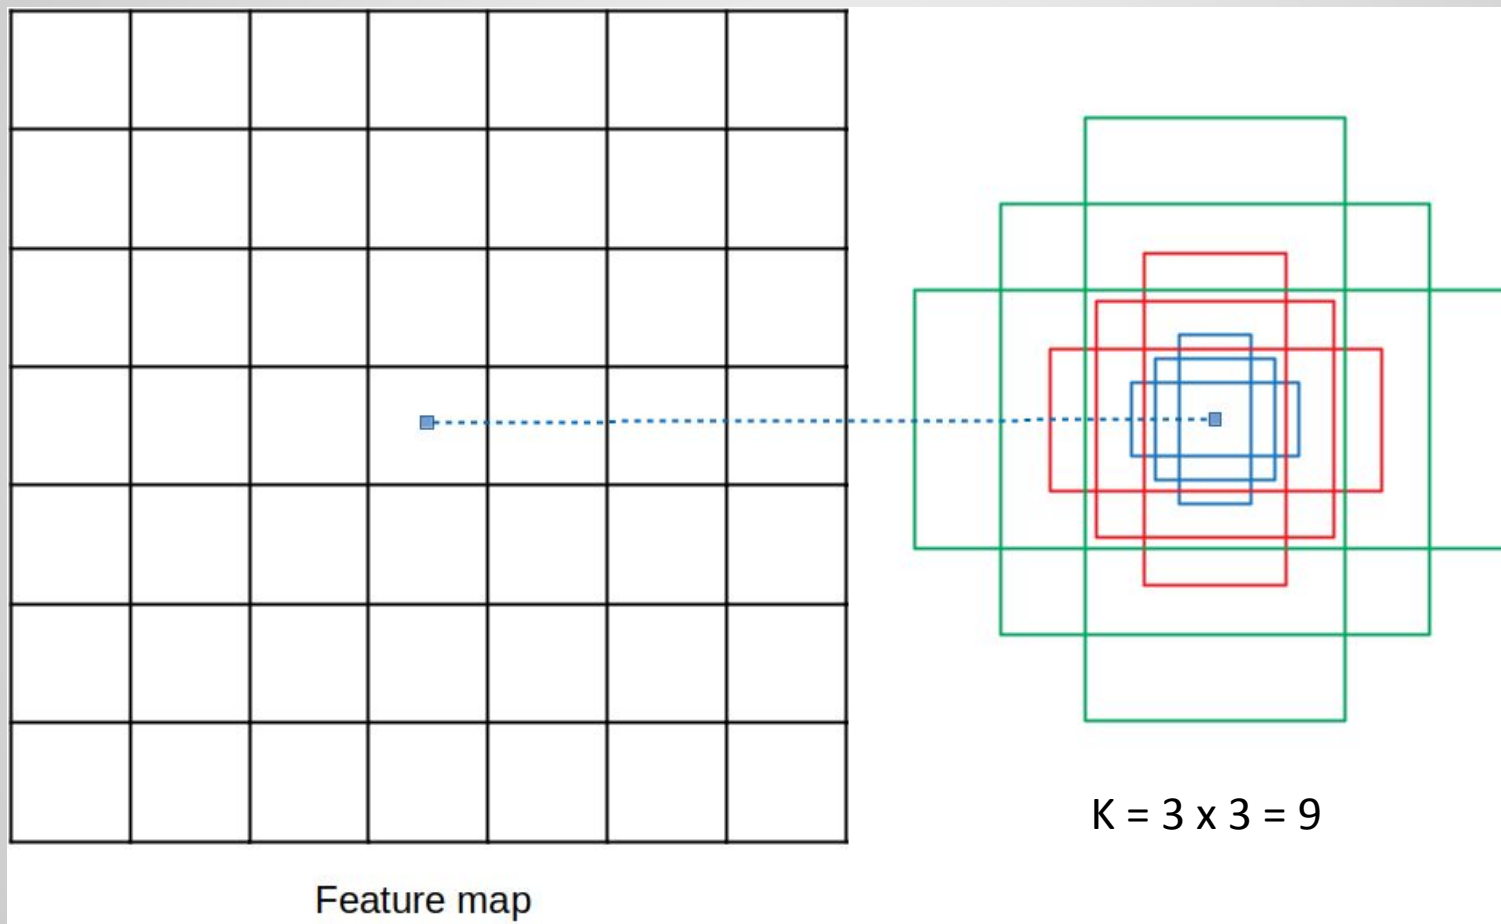
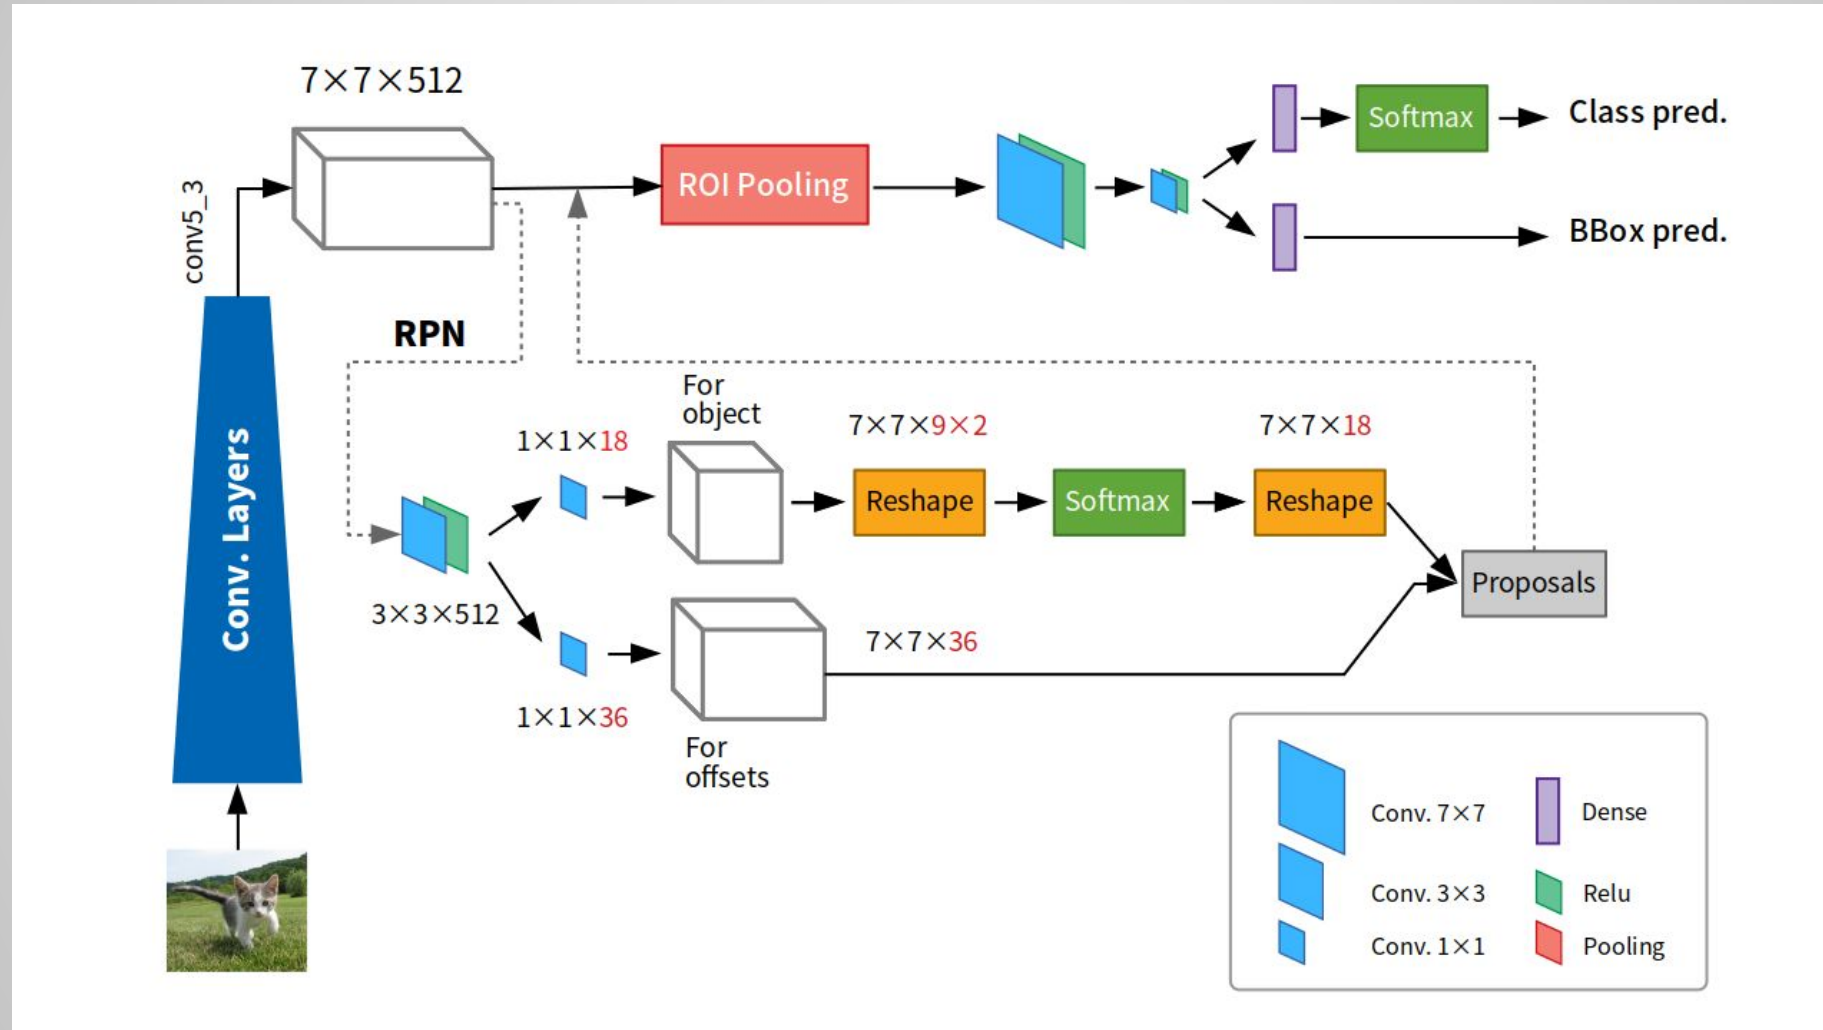


Figure 2: Faster R-CNN is a single, unified network for object detection. The RPN module serves as the 'attention' of this unified network.

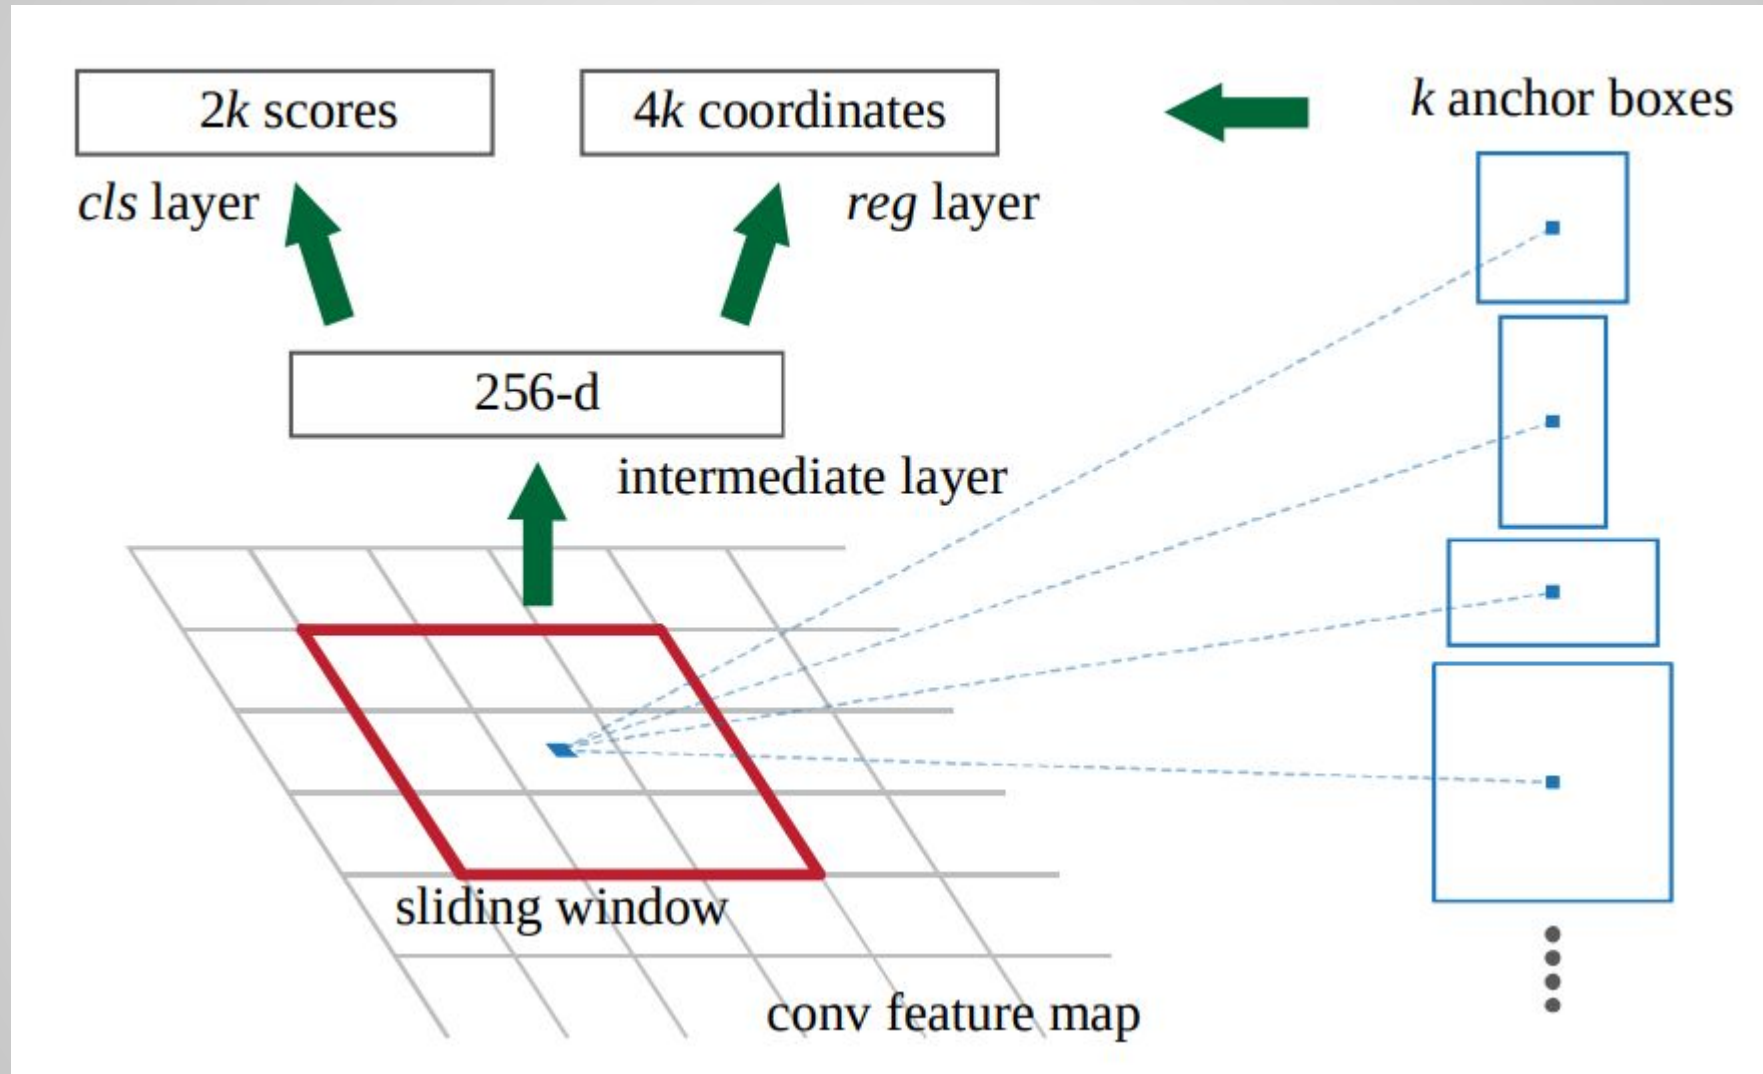
Anchor




Region Proposal Network (RPN)



Region Proposal Network (RPN)



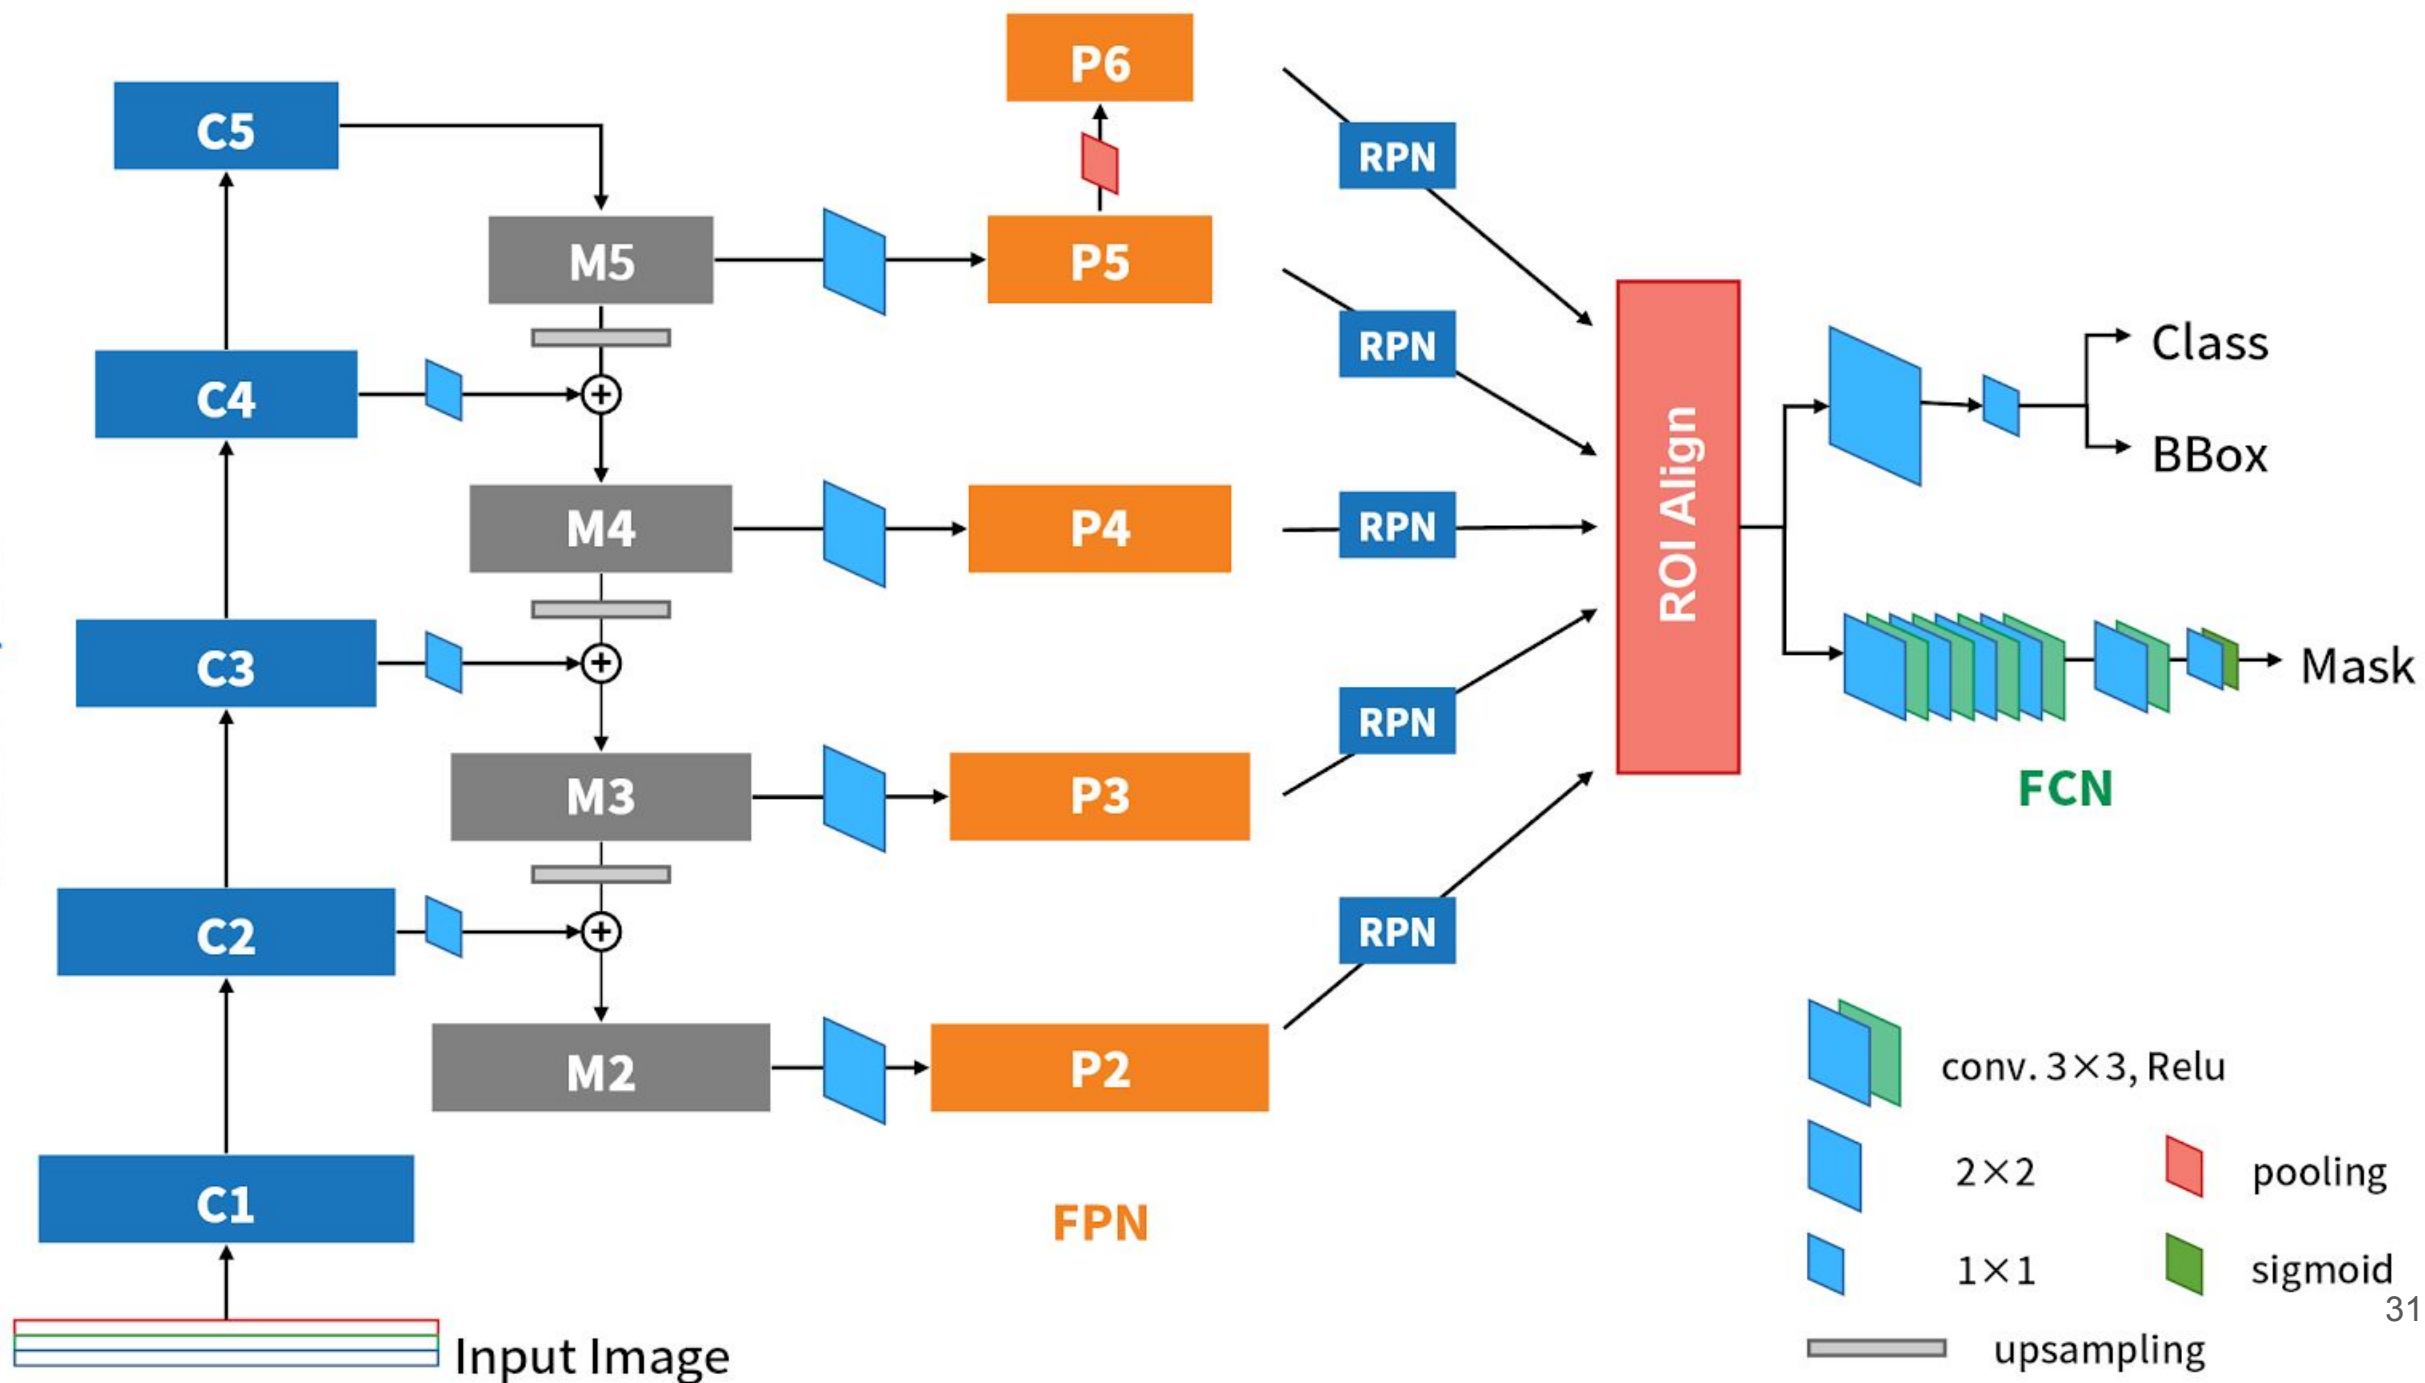
Region Proposal Network (RPN)




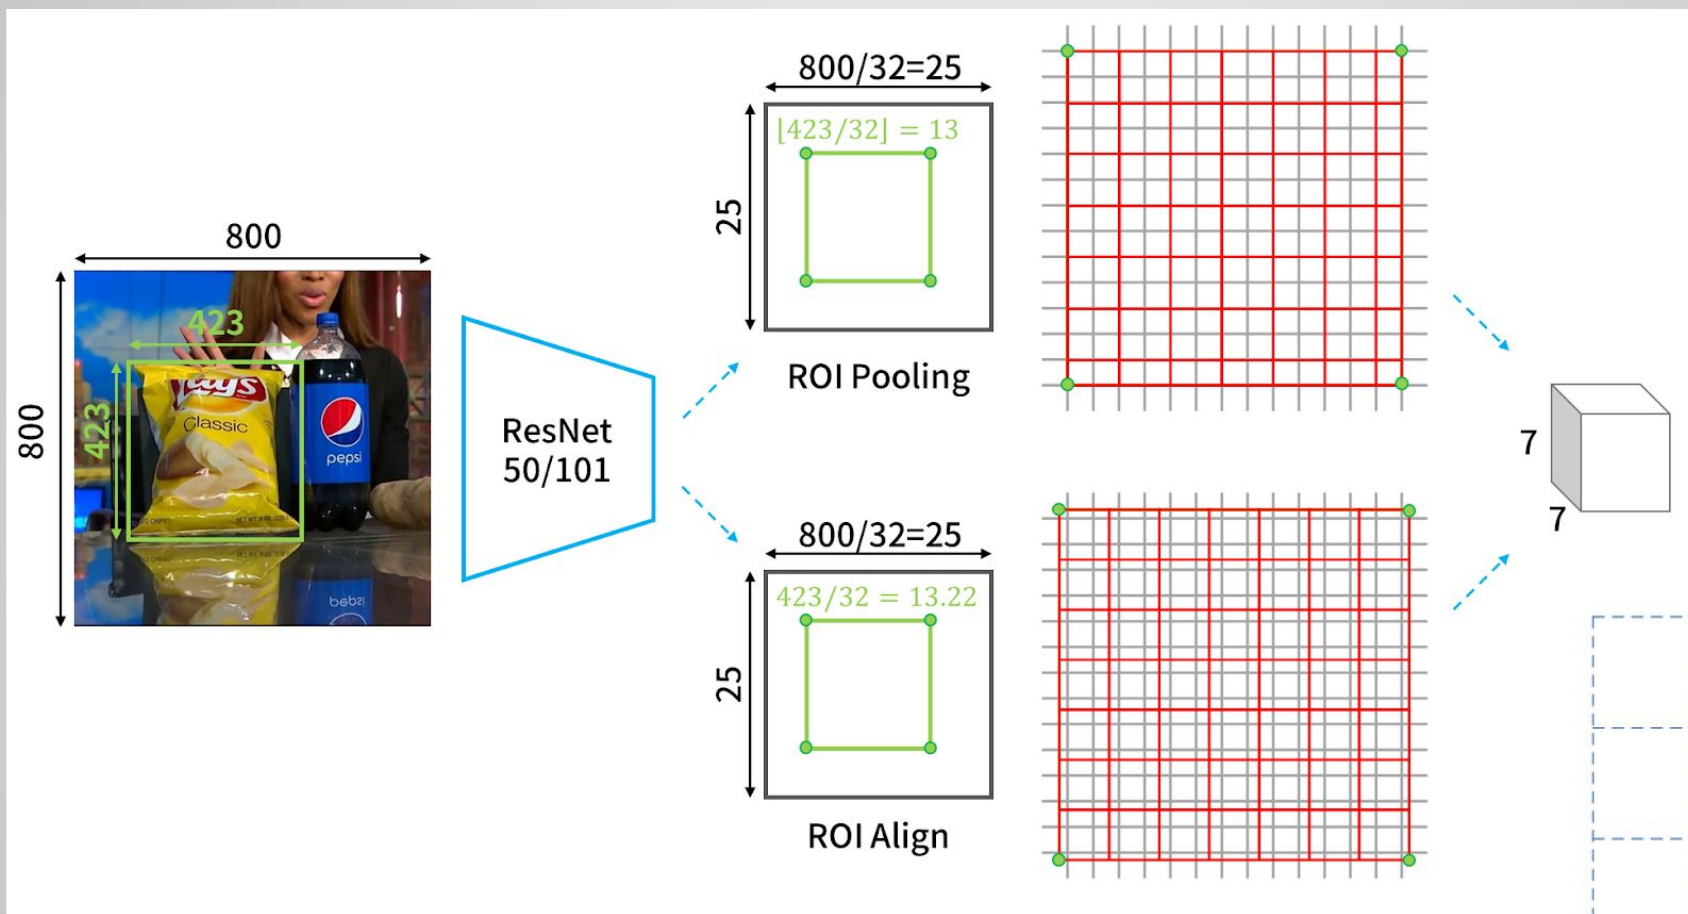
RPN Loss Function

$$\begin{aligned} L(\{p_i\}, \{t_i\}) = & \frac{1}{N_{cls}} \sum_i L_{cls}(p_i, p_i^*) \\ & + \lambda \frac{1}{N_{reg}} \sum_i p_i^* L_{reg}(t_i, t_i^*). \end{aligned} \quad (1)$$

Mask R-CNN




ROI Align



Detection

Tomato


Label



Detection

Tomato

Predict




Detection


Tomato

fast_rcnn 3 cards


fast_rcnn/cls_accuracy



fast_rcnn/false_negative



fast_rcnn/fg_cls_accuracy



Code :

https://colab.research.google.com/drive/1uy_55Ziq3U-FBIOlv-v8A2Hy3od4rWyr?usp=sharing

Reference :

<https://github.com/facebookresearch/detectron2>
<https://arxiv.org/abs/1703.06870>

For more information about Tomato Detection Demo, ask us in QA time :

<https://hackmd.io/@2peEy1j8QSygBgdKIEZpJQ/HkIsfQih2>